



Tool benchmark -- RUNNING comparison (2 external tools in)


Precision vs PolyASite 2.0 rep sites (point mode, strand-matched)




Recall vs atlas in 14,851 detected-gene bodies (shared denominator)




Sites called



Runtime (not comparable: scope differs)



Peak memory



- polyApipe (misprime excl.)
- polyApipe (all peaks)
- Sierra (3' of fit)
- Sierra (summit)
- PeakATail lg_annotate
- PeakATail B1_cohort_full
- ◆ genic-shuffle null (3 seeds)
- * provisional (Laughney cohort)

CAVEATS: PeakATail bars (*) are PROVISIONAL -- scored on the Laughney cohort, not pbmc_10k_v3; the matched CellRanger pbmc_10k_v3 PeakATail run is still in progress. polyApipe and Sierra scored on pbmc_10k_v3 (CellRanger BAM); Sierra reduced to the strand-aware 3' base of Fit.start..Fit.end (primary) or MaxPosition summit (sensitivity). Cross-dataset precision/recall differences partly reflect dataset depth/composition, not tool quality alone. Runtime/RSS scope differs: polyApipe/Sierra = PAS calling + per-cell counting (/usr/bin/time, one run for both variants); PeakATail = full pipeline incl. clustering/differential (5 s process sampling, may miss subprocess peaks). Recall denominator identical for all tools: 285,220 PolyASite 2.0 rep sites in the 14,851 Laughney-detected genes (recall flavor b of benchmark_curated.py).