



● include_extended = on (distance sweep)


◇ include_extended = off (pipeline default)



band = between-dataset IQR, ~6× wider than the entire distance effect · n = 17 datasets · at the pipeline default, 23.1% of called peaks are retained (16.8–28.6%); the filter thresholds that set this were NOT swept




band = between-dataset IQR · median is flat to the pixel: worst single dataset moves 3.8% over all 13 branches, and no branch has a median shift above 0.14% · n = 17 datasets



clustering branch — all six share one identical preprocessed matrix

bars = median, whiskers = IQR across 17 datasets · counted as unique obs['leiden'] labels in 07_clustering/<dataset>/clusters.h5ad
cluster COUNT is a granularity proxy only — equal counts do not imply equal cell assignments; partition agreement is not tested here (see clustering_concordance)



bars = median per-dataset paired change between the two ends of that knob, whiskers = IQR over 17 datasets; the n/17 tag counts datasets that changed at all. The two arms are NOT symmetric: clustering knobs leave PAS and cell counts bit-identical (0/17 datasets move, exactly 0.00%), whereas annotation knobs leave the MEDIAN cluster count unchanged but do shift 7–9 of 17 individual datasets by 1–2 clusters (–13% to +13%) — a median-zero bar is not a null result. Also: tfidf→libsize has a small median (+6%) but the widest per-dataset spread of any knob (–42% to +93%), so it is stable only on average; and the filter thresholds (min_read/min_cells/min_pas_per_cell), plausibly the larger lever on PAS counts, are held fixed in all 13 branches and are not tested here.